

Asset Defender Risk Logic Q&A Summary

Generated on 2026-03-07

Q&A Summary

1) Source of truth in production

- Live popup verdict: popup-computed finalRisk.
- Persisted/reporting history: server final_predictions.risk_score.
- Current state: not guaranteed identical.

2) Unify weighted formula (0.3/0.7 vs 0.6/0.4)

- Yes, unify everywhere.
- Use one shared formula constant across popup, server, and reporting UI.

3) Rename “Moderate” / “modern” wording

- Yes, rename for clarity.
- Recommended labels: Very Low, Low, Medium, High, Critical.

4) Hard-threshold docs in one table

- Yes.
- Include risk bands, category rules (hacker/scam/bot/suspicious/genuine), override rules, and boundary examples.

5) Owner of final classification logic

- Recommended production owner: server.
- Popup can show provisional output until server confirmation.
- If both paths remain, enforce parity tests.

6) Persist final category + exact reasons

- Yes, required for auditability.
- Persist category, score, triggered rule IDs, feature values, model/heuristic version, timestamp, and source.

7) Conflict policy when popup and server disagree

- Server wins for stored/reporting/enforcement decisions.
- UI shows provisional client result, then replaces with confirmed server result.
- Log mismatch diffs.

8) Regression tests for thresholds (39/40, 69/70, etc.)

- Yes.
- Add below/at/above boundary tests, override-path tests, insufficient-evidence tests, and parity tests.

Additional Questions Recommended

1. Should popup scoring be production-disabled and used only for offline/debug mode?
2. Should every decision store a deterministic decision-trace payload for replay?
3. What server confirmation SLA should replace provisional client output?
4. Which threshold set is versioned as v1.0, and how do we roll back v1.1 safely?